

KRISHNA GUNDIMEDA

+(91) 70322-97414 ◊ Hyderabad, India

gkrishnaprabhat@gmail.com ◊ linkedin.com/in/krishna-prabhat-gundimeda ◊ kris225.github.io

Frontend and AI Engineer with experience building production-grade web applications and intelligent systems, with a strong foundation in Data Structures and Algorithms.

EDUCATION

Bachelor of Technology in Computer Science, Vellore Institute of Technology, Vellore
GPA: 8.16 / 10

2020 - 2024

SKILLS

Frontend	JavaScript (ES6+), React, AngularJS, StencilJS, HTML, CSS, Tailwind CSS
AI / ML	Python, Vertex AI, RAG/CAG, MCP, Google ADK, LLM Tool Use
Backend	Node.js, Python, SQL, REST APIs, ETL Pipelines, Pandas
Tools	Git, VS Code, Jupyter, GCP, Jaspersoft, Figma

EXPERIENCE

Frontend & AI Engineer @ Broadcom
Futurist Global

Sep 2024 - Present
Hyderabad, India

- Enabled intelligent querying and summarization across 500K+ enterprise records by shipping MCP-based services for Broadcom's ICE and Clarity products.
- Accelerated migration of the front-end codebase primarily written in Angular 1.x to Stencil JS, a web component based framework by converting 16+ reusable components and workflows.
- Improved data retrieval accuracy and reduced stakeholder turnaround time by implementing internal AI tools powered by Vertex AI, RAG pipelines, and A2A architectures for automated report generation and contextual insights.

STEM Intern @ Goldman Sachs
Goldman Sachs

Jan 2024 - Jun 2024
Bangalore, India

- Built Python-based data analysis and ETL workflows for middle-office CD operations.
- Automated reporting pipelines, saving over 195+ hours per day across internal teams.
- Worked with large datasets and internal tooling to improve reliability and turnaround time of analytics jobs.

PERSONAL PROJECTS AND HACKATHONS

Reportly for Clarity PPM

RAG, Google ADK, Vertex AI, Python, React

- Project won Broadcom's ValueOps AI Hackathon.
- Built an AI-powered reporting tool that automates Jaspersoft report generation using agent-based workflows.
- Uses intelligent agents to understand user queries, retrieve relevant enterprise data across large datasets, to generate UI of the report dynamically.

CAG - Cache-Augmented Generation

Python, Vertex AI, Embeddings

- Built a lightweight alternative to RAG that grounds LLM responses using custom documentation with semantic chunking and vector embeddings via text-embedding-005.
- Features cosine similarity search for context retrieval, chat history summarization to reduce token costs, and token usage tracking with cost estimation.

Binary Star System Simulation

Python, NumPy, Matplotlib, ODE Solvers

- Project won VIT Vellore's Gravathon 2023. Built a simulation of planetary motion in a binary star system, solving the gravitational three-body problem using numerical methods.
- Implemented Runge-Kutta integration for accurate orbital mechanics, computing stable and chaotic trajectories influenced by two gravitationally bound stars.

BLOGS

- [The Case for CAG: Less Retrieval, More Speed](#) - Exploring Cache-Augmented Generation as an alternative to RAG
- [Gemini: Google's Multimodal Moonshot](#) - Technical analysis of Google's Gemini AI model family
- [Llama 2: How Meta Gave Away the Recipe \(Mostly\)](#) - Deep dive into Meta's Llama 2 release and training methodology

CERTIFICATIONS

- Machine Learning Specialization - DeepLearning.AI & Stanford University (Andrew Ng)
- The Web Developer Bootcamp - Udemy
- 100 Days of Code: Python - Udemy
- Google Cloud Skills Boost - Google Cloud